

# AI for Genomics: from CNNs and LSTMs to Transformers

Nikolay Oskolkov, Group Leader (PI) at LIOS, Riga, Latvia

Physalia course, 09.09.2025

## Session 2a: Convolutional Neural Networks (CNNs) applications for Microscopy Imaging



@NikolayOskolkov



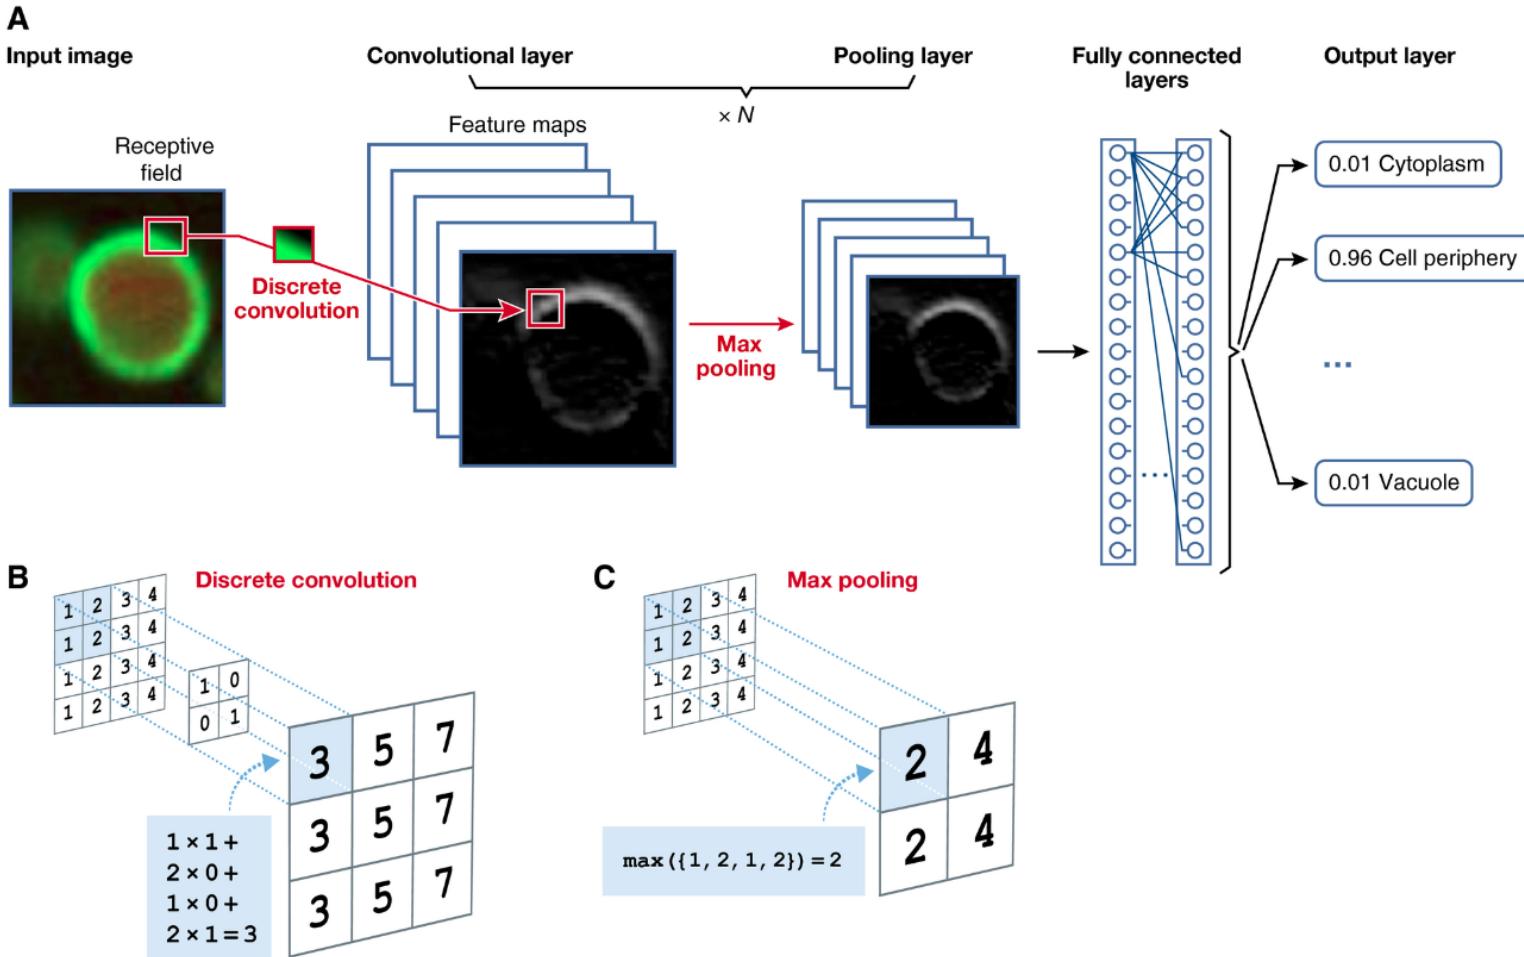
@oskolkov.bsky.social



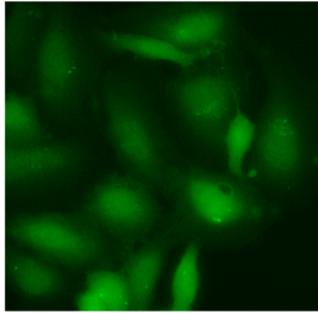
Personal homepage:  
<https://nikolay-oskolkov.com>

Topics we'll cover in this session:

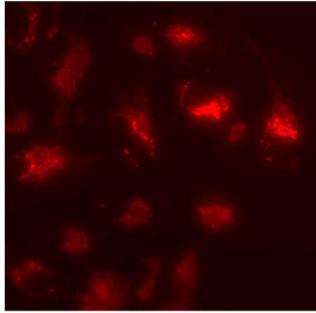
- 1) Introduction to Convolutional Neural Networks (CNNs)
- 2) Microscopy imaging: object detection and segmentation
- 3) Grad-CAM for interpreting Deep Learning models
- 4) Unsupervised clustering of images and batch detection
- 5) Comparing Faster RCNN and Mask-RCNN for object detection



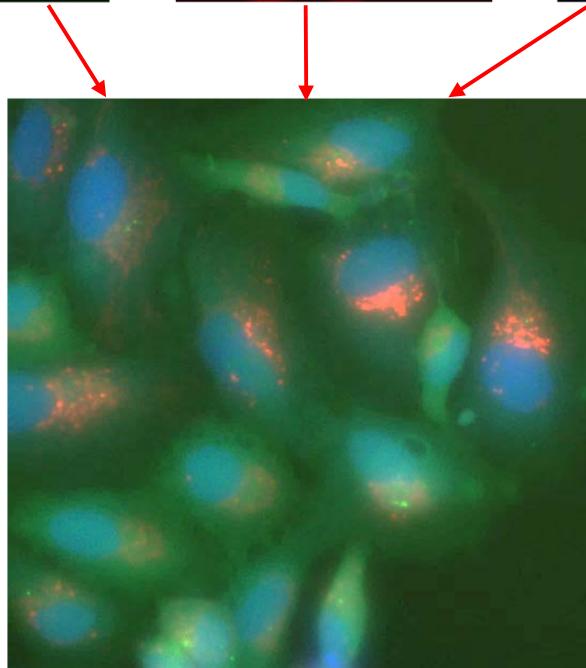
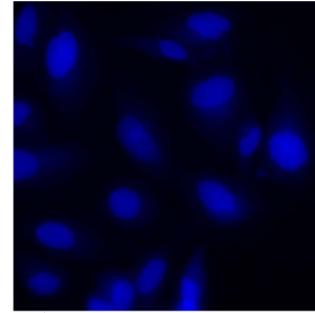
Cytoplasm



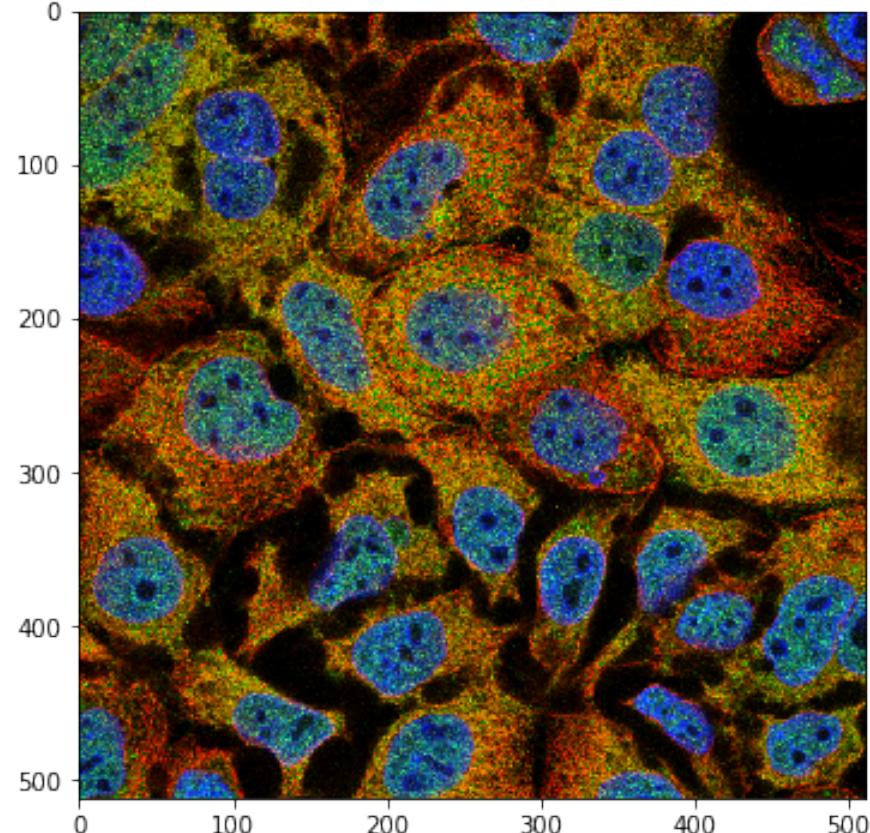
Mitochondria



Nucleus

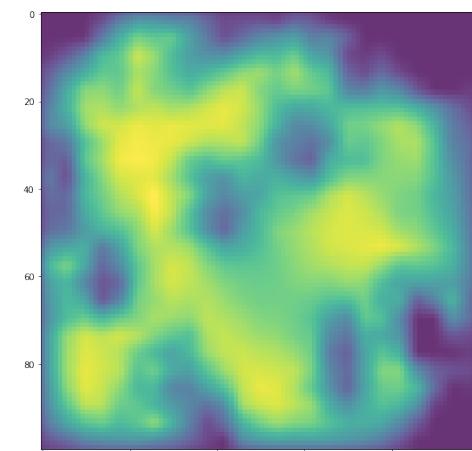
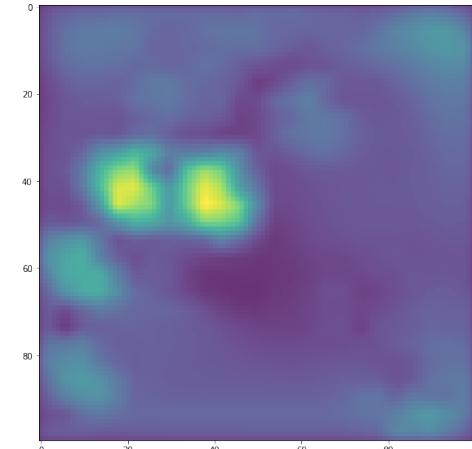
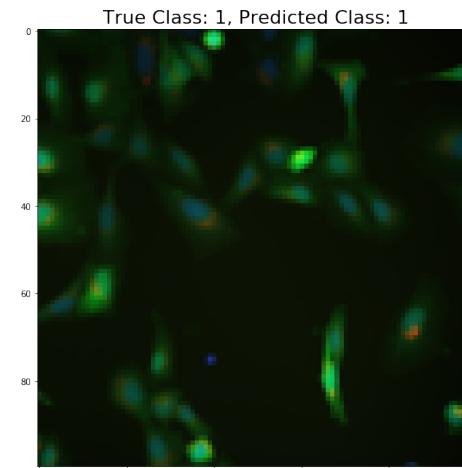
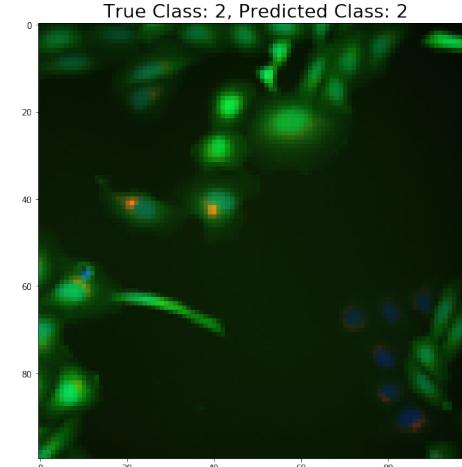
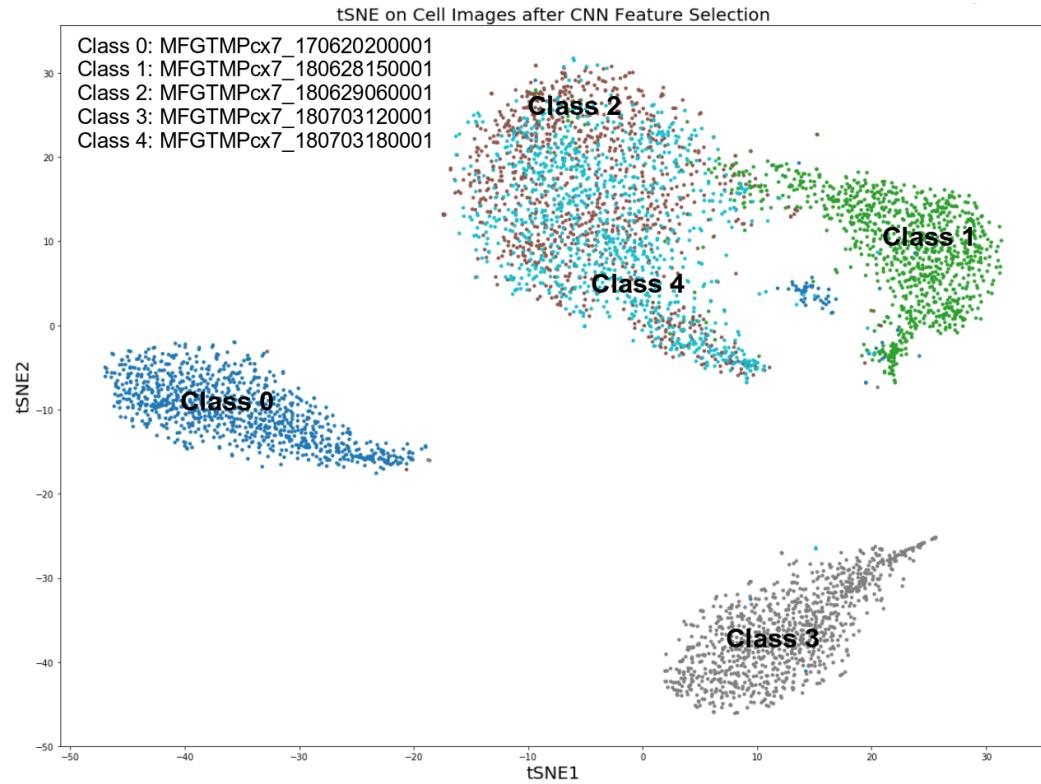


Human Protein Atlas data



One way to merge the 4 channels in Python would be to realize that yellow = red + green and add half of the yellow channel to the red and the other half to the green channel.

# Grad-CAM to interpret image clusters



Different batches had different background

## OpenCV-Python Tutorials

### Rotation

Rotation of an image for an angle  $\theta$  is achieved by the transformation matrix of the form

$$M = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix}$$

But OpenCV provides scaled rotation with adjustable center of rotation so that you can rotate at any location you prefer. The modified transformation matrix is given by

$$\begin{bmatrix} \alpha & \beta & (1-\alpha) \cdot \text{center}_x - \beta \cdot \text{center}_y \\ -\beta & \alpha & \beta \cdot \text{center}_x + (1-\alpha) \cdot \text{center}_y \end{bmatrix}$$

where:

$$\begin{aligned} \alpha &= \text{scale} \cdot \cos\theta, \\ \beta &= \text{scale} \cdot \sin\theta \end{aligned}$$

To find this transformation matrix, OpenCV provides a function, `cv.getRotationMatrix2D`. Check out the below example which rotates the image by 90 degrees with respect to center without scaling.

```
img = cv.imread('messi5.jpg', cv.IMREAD_GRAYSCALE)
assert img is not None, "file could not be read, check with os.path.exists()"
rows,cols = img.shape
```

`M` col-1 and rows-1 are the coordinate limits.

`M = cv.getRotationMatrix2D([cols-1/2.0, [rows-1/2.0], 90])`

See the result:



### Canny Edge Detection in OpenCV

OpenCV puts all the above in single function, `cv.Canny()`. We will see how to use it. First argument is our input image. Second and third arguments are `minVal` and `maxVal` respectively. Fourth argument is `aperture_size`. It is the size of Sobel kernel used for finding gradients. By default it is 3. Last argument is `Laplacian` which specifies the equation for finding gradient magnitude. If it is True, it uses the equation mentioned above which is more accurate, otherwise it uses this function:  $\text{Edge Gradient } (G_e) = |G_x| + |G_y|$ . By default, it is False.

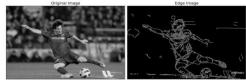
```
import numpy as np
import cv2
from matplotlib import pyplot as plt

img = cv2.imread('messi5.jpg', cv.IMREAD_GRAYSCALE)
assert img is not None, "file could not be read, check with os.path.exists()"
edges = cv2.Canny(img, 100, 200)

plt.subplot(121),plt.imshow(edges,cmap = 'gray')
plt.title('Edge Image'), plt.xticks([]), plt.yticks([])
plt.subplot(122),plt.imshow(edges,cmap = 'gray')
plt.title('Edge Image'), plt.xticks([]), plt.yticks([])

plt.show()
```

See the result below:



### Introduction to OpenCV

Learn how to setup OpenCV-Python on your computer!

### Gui Features in OpenCV

Here you will learn how to display and save images and videos, control mouse events and create trackbar.

### Core Operations

In this section you will learn basic operations on image like pixel editing, geometric transformations, code optimization,

### Image Processing in OpenCV

In this section you will learn different image processing functions inside OpenCV.

### Feature Detection and Description

In this section you will learn about feature detectors and descriptors

### Video analysis (video module)

In this section you will learn different techniques to work with videos like object tracking etc.

### Camera Calibration and 3D Reconstruction

In this section we will learn about camera calibration, stereo imaging etc.

### Machine Learning

In this section you will learn different image processing functions inside OpenCV.

### Computational Photography

In this section you will learn different computational photography techniques like image denoising etc.

### Object Detection (objdetect module)

In this section you will learn object detection techniques like face detection etc.

### OpenCV-Python Bindings

In this section, we will see how OpenCV-Python bindings are generated

```
import numpy as np
from matplotlib import pyplot as plt

img = cv.imread('noisy2.png', cv.IMREAD_GRAYSCALE)
assert img is not None, "file could not be read, check with os.path.exists()"

# global thresholding
ret1,th1 = cv.threshold(img,127,255,cv.THRESH_BINARY)

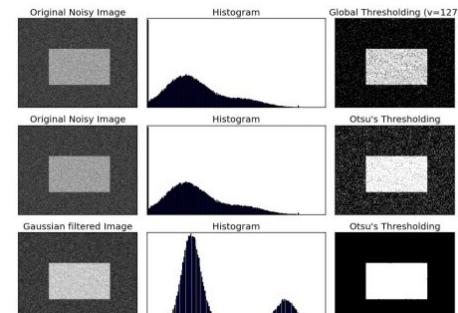
# Otsu's thresholding
ret2,th2 = cv.threshold(img,0,255,cv.THRESH_BINARY+cv.THRESH_OTSU)

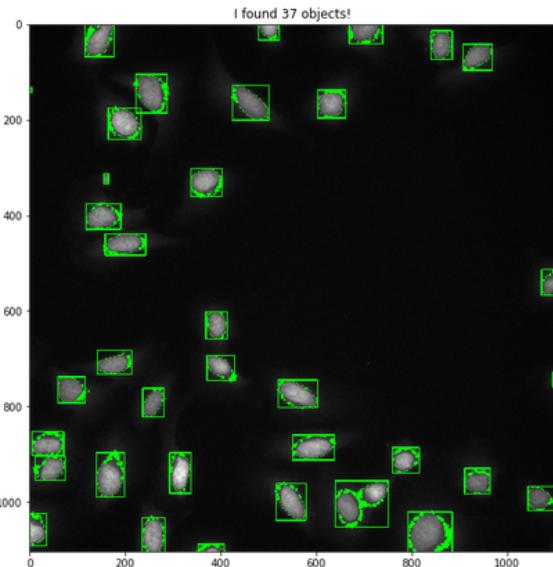
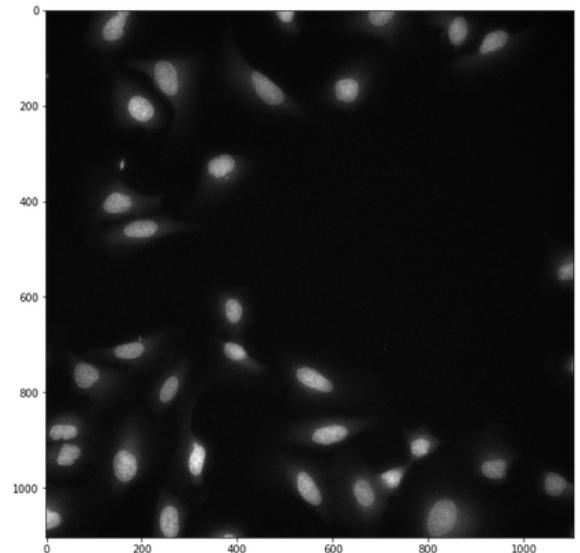
# Otsu's thresholding after Gaussian filtering
blur = cv.GaussianBlur(img,(5,5),0)
ret3,th3 = cv.threshold(blur,0,255,cv.THRESH_BINARY+cv.THRESH_OTSU)

# plot all the images and their histograms
images = [img, th1,
          img, th2,
          blur, th3]
titles = ['Original Noisy Image','Histogram','Global Thresholding (v=127)', 
          'Original Noisy Image','Histogram','Otsu's Thresholding',
          'Gaussian filtered Image','Histogram','Otsu's Thresholding']

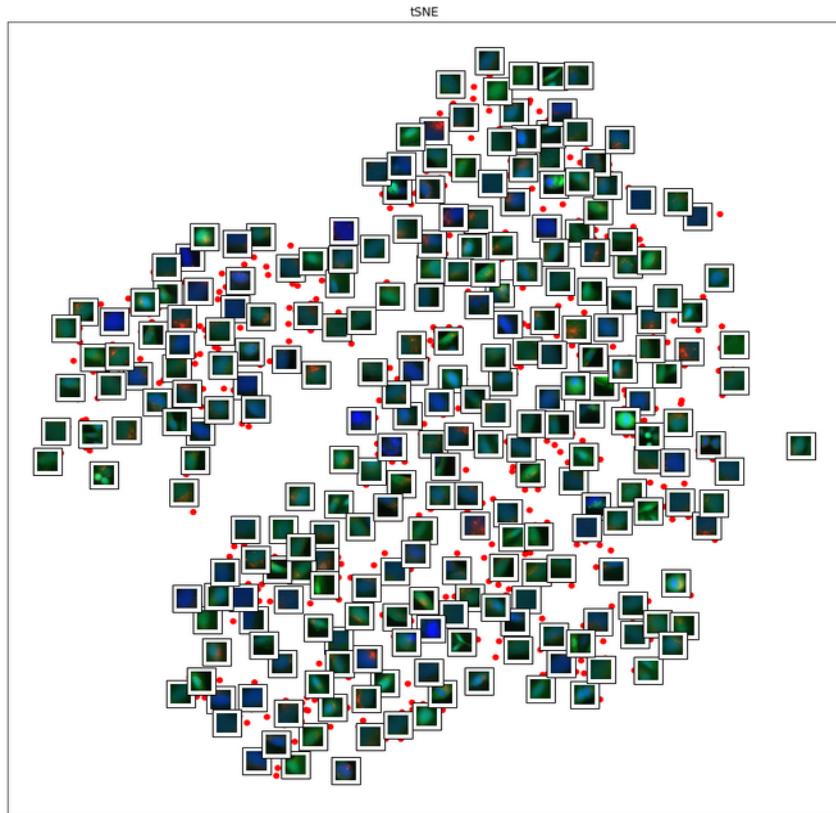
for i in range(3):
    plt.subplot(3,3,i*3+1),plt.imshow(images[i*3],'gray')
    plt.title(titles[i*3]), plt.xticks([]), plt.yticks([])
    plt.subplot(3,3,i*3+2),plt.hist(images[i*3].ravel(),256)
    plt.title(titles[i*3+1]), plt.xticks([]), plt.yticks([])
    plt.subplot(3,3,i*3+3),plt.imshow(images[i*3+2],'gray')
    plt.title(titles[i*3+2]), plt.xticks([]), plt.yticks([])
plt.show()
```

Result:

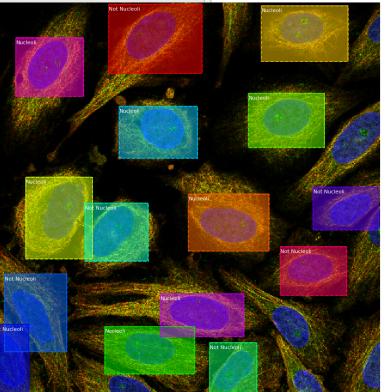
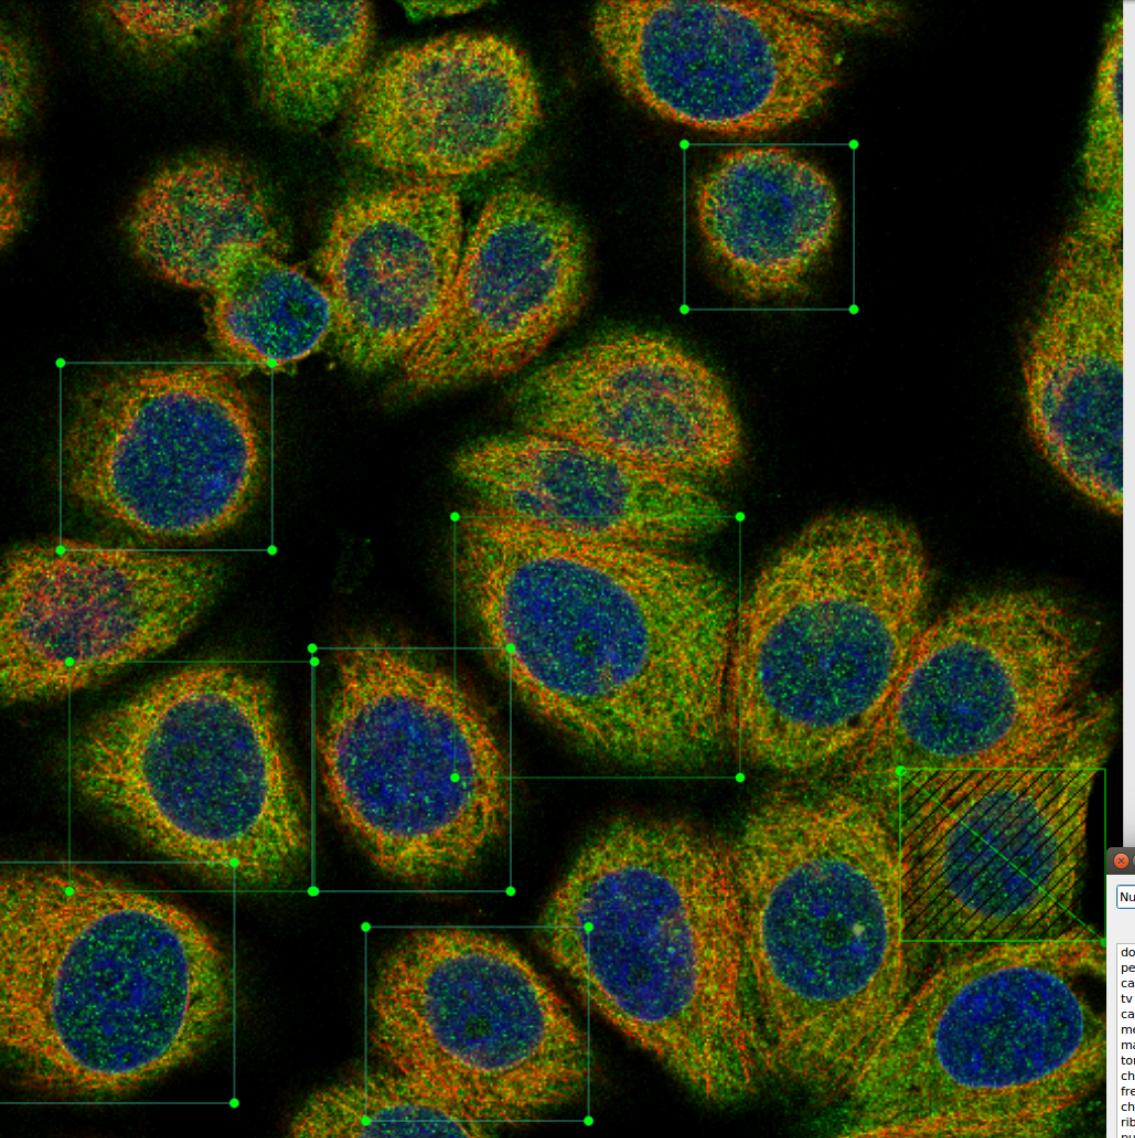




Cells were located by thresholding and cropped.

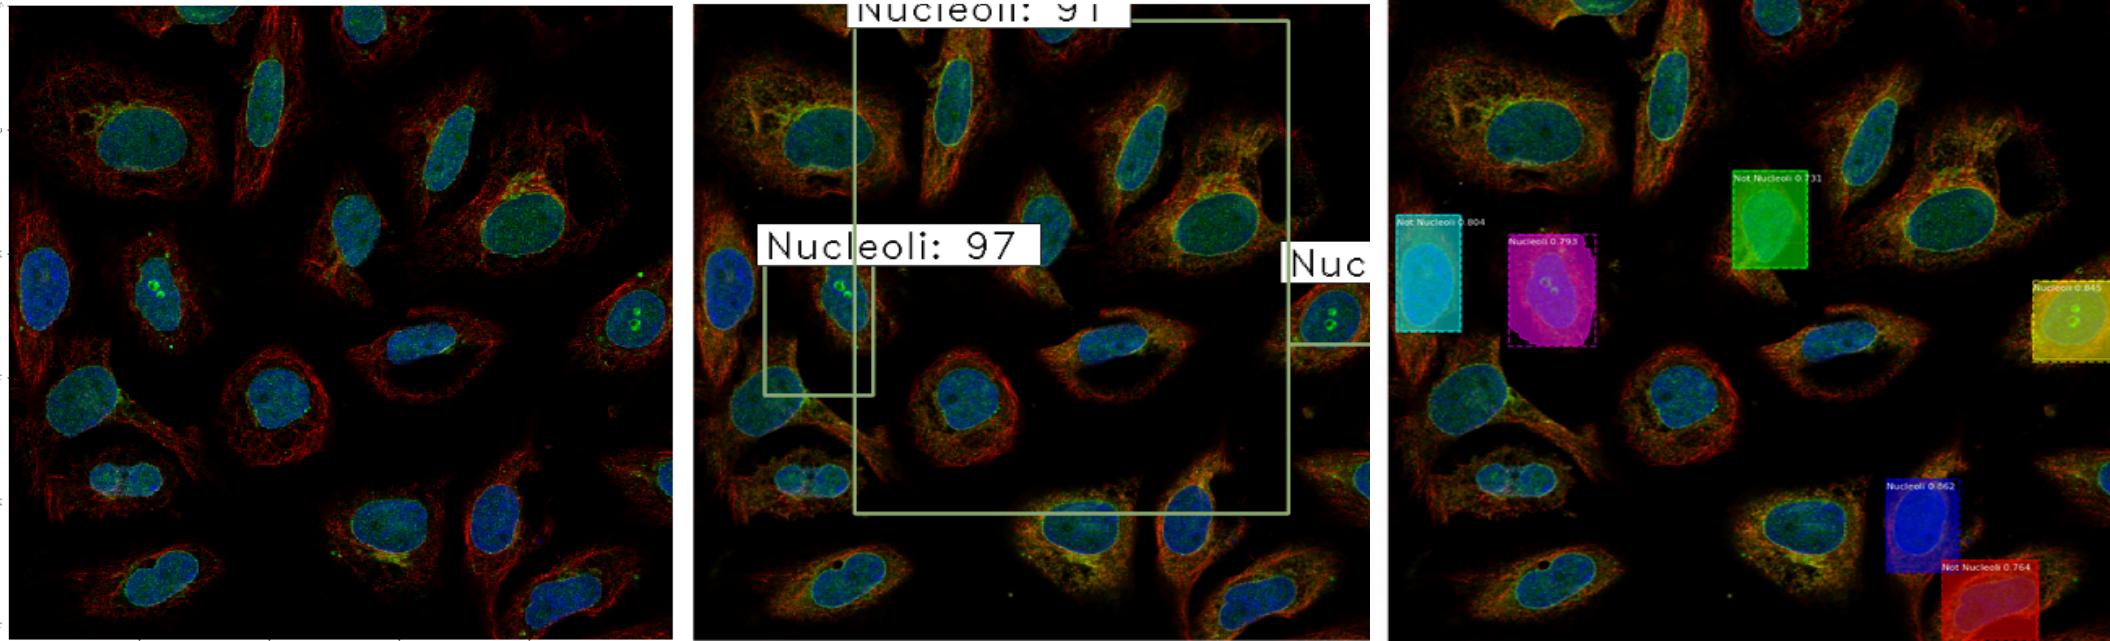


Features were extracted by Convolutional Autoencoder.  
Cells were visualized by tSNE using the bottleneck  
layer of Convolutional Autoencoder.



A screenshot of a Linux desktop environment showing a "labeling" dialog box. The title bar says "labeling". Inside, there is a text input field containing "NucPlusCyt" with a cursor. Below the input field are two buttons: a red "Cancel" button with a left-pointing arrow icon and a green "OK" button with a checkmark icon. A scroll bar is visible on the right side of the dialog box. The background shows other windows and icons typical of a Linux desktop.

# Faster RCNN vs. Mask-RCNN for Nucleoli object detection



Faster-RCNN had lower accuracy of cell detection and was much slower than Mask-RCNN

## Cell Detection with Star-convex Polygons

Uwe Schmidt<sup>1,\*</sup>, Martin Weigert<sup>1,\*</sup>, Coleman Broaddus<sup>1</sup>, and Gene Myers<sup>1,2</sup>

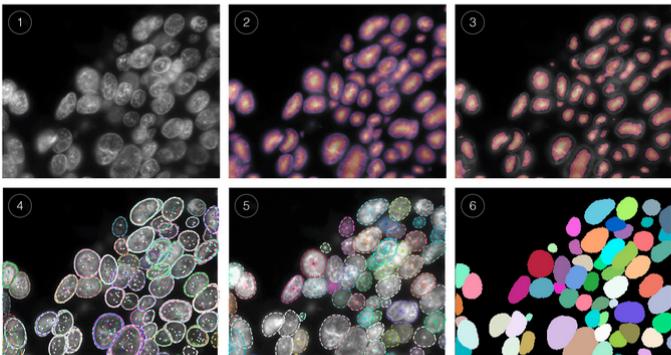
<sup>1</sup> Max Planck Institute of Molecular Cell Biology and Genetics, Dresden, Germany  
Center for Systems Biology Dresden, Germany

<sup>2</sup> Faculty of Computer Science, Technical University Dresden, Germany

**Abstract.** Automatic detection and segmentation of cells and nuclei in microscopy images is important for many biological applications. Recent successful learning-based approaches include per-pixel cell segmentation with subsequent pixel grouping, or localization of bounding boxes with subsequent shape refinement. In situations of crowded cells, these can be prone to segmentation errors, such as falsely merging bordering cells or suppressing valid cell instances due to the poor approximation with bounding boxes. To overcome these issues, we propose to localize cell nuclei via *star-convex polygons*, which are a much better shape representation as compared to bounding boxes and thus do not need shape refinement. To that end, we train a convolutional neural network that predicts for every pixel a polygon for the cell instance at that position. We demonstrate the merits of our approach on two synthetic datasets and one challenging dataset of diverse fluorescence microscopy images.

### Usage

We provide example workflows for 2D and 3D via Jupyter notebooks that illustrate how this package can be used.



Pretrained Models for 2D

## Cellpose: a generalist algorithm for cellular segmentation

Carsten Stringer<sup>1,\*</sup>, Michalis Michaelos<sup>1</sup>, Marius Pachitariu<sup>1\*</sup>

<sup>1</sup>HHMI Janelia Research Campus, Ashburn, VA, USA

\* correspondence to (stringerc, pachitarium) @ janelia.hhmi.org

Many biological applications require the segmentation of cell bodies, membranes and nuclei from microscopy images. Deep learning has enabled great progress on this problem, but current methods are specialized for images that have large training datasets. Here we introduce a generalist, deep learning-based segmentation algorithm called Cellpose, which can very precisely segment a wide range of image types out-of-the-box and does not require model retraining or parameter adjustments. We trained Cellpose on a new dataset of highly-varied images of cells, containing over 70,000 segmented objects. To support community contributions to the training data, we developed software for manual labelling and for curation of the automated results, with optional direct upload to our data repository. Periodically retraining the model on the community-contributed data will ensure that Cellpose improves constantly.

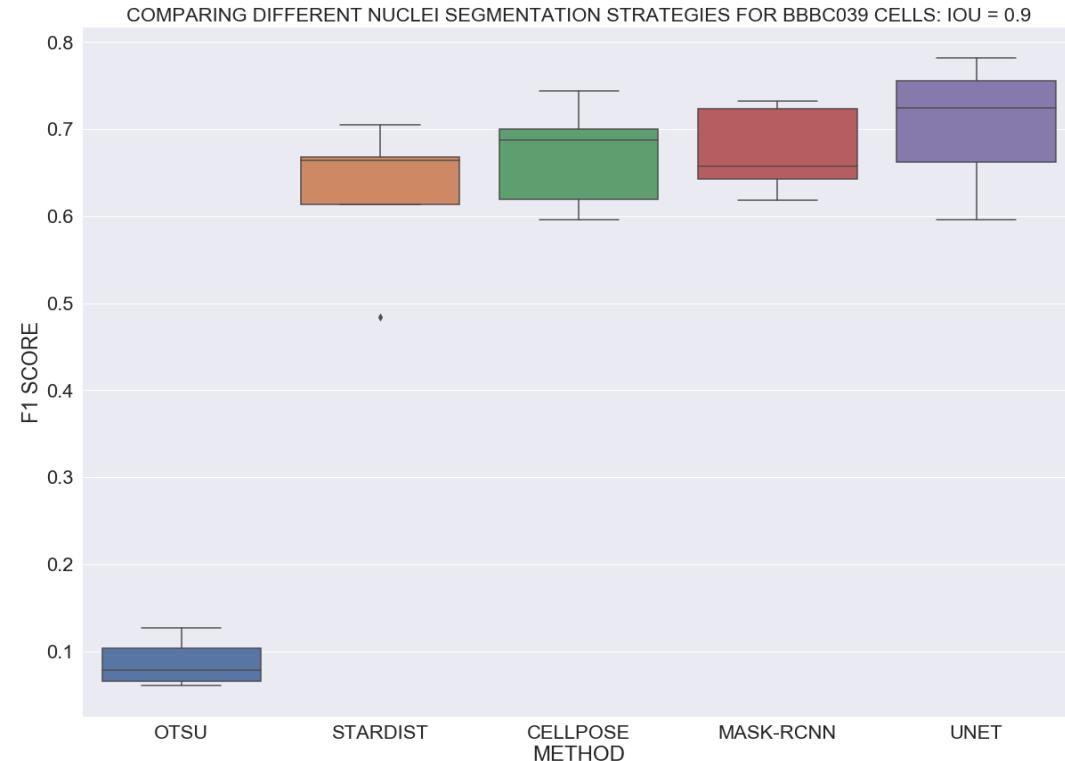
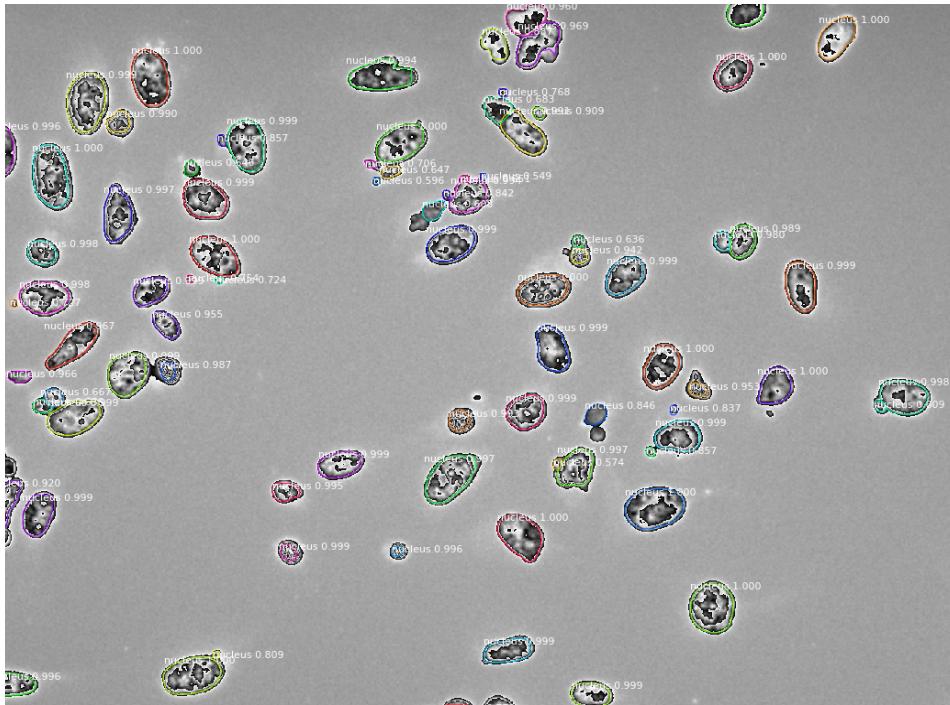
### Introduction

Quantitative cell biology requires simultaneous measurements of different cellular properties, such as shape, position, RNA expression and protein expression [1]. A first step towards assigning these properties to single cells is the segmentation of an imaged volume into cell bodies, usually based on a cytoplasmic or membrane marker [2–8]. This step can be straightfor-

puter vision algorithms like Mask R-CNN [15, 16] and adapted the algorithms for the biological problem. Following the competition, this dataset generated further progress, with other methods like Stardist and nucle-Alzer being developed specifically for this dataset [17, 18].

In this work, we followed the approach of the Data Science Bowl team to collect and segment a large dataset of cell images from a variety of microscop

The screenshot shows the Cellpose web application interface. At the top, there is a header with the Cellpose logo and a brief description: "a generalist algorithm for cellular segmentation". Below the header, there is a section with a dashed green box for file upload and a note: "Drop files here or click to upload." and "Try cellpose by uploading one PNG or JPG <10 MB. Images are resized to a max size of 512x512 pixels." Below this, there is a section titled "or click on an example image from our test set:" followed by a grid of small microscopy images showing various cell types and patterns.



Strategy: use UNET / Mask-RCNN for segmenting and cropping cells, and run clustering

[README](#) [Code of conduct](#) [AGPL-3.0 license](#) [Security](#)

ultralytics  
YOLO AI  
Unleashing Next-Gen AI

Discover more

- Enhanced feature extraction architecture
- Faster processing with balanced accuracy
- Higher precision using 22% fewer parameters
- Easily deployable on edge, cloud, and GPU systems
- YOLO11 enterprise models providing enterprise-grade accuracy

Download the App

Available on the [App Store](#) [GET IT ON Google Play](#)

中文 | 한국어 | 日本語 | Русский | Deutsch | Français | Español | Português | Türkçe | Tiếng Việt | العربية

Ultralytics CI passing | downloads 97M | DOI 10.5281/zenodo.7347926 | Discord 819 online | Forums 419 users | Reddit 568 | Run on Gradient | Open in Colab | Open in Kaggle | launch binder

ultralytics creates cutting-edge, state-of-the-art (SOTA) [YOLO models](#) built on years of foundational research in computer vision and AI. Constantly updated for performance and flexibility, our models are **fast**, **accurate**, and **easy to use**. They excel at [object detection](#), [tracking](#), [instance segmentation](#), [image classification](#), and [pose estimation](#) tasks.

Find detailed documentation in the [Ultralytics Docs](#). Get support via [GitHub Issues](#). Join discussions on [Discord](#), [Reddit](#), and the [Ultralytics Community Forums](#)!

Request an Enterprise License for commercial use at [Ultralytics Licensing](#).

Model	Latency (ms/img)	COCO mAP50-95 (%)
YOLO11	~1.5	~40.5
YOLO10	~2.5	~45.5
YOLO9	~3.5	~48.5
YOLO8	~4.5	~49.5
YOLO7	~5.5	~50.5
YOLO6-3.0	~6.5	~51.5
YOLO5	~7.5	~52.5
PP-YOLOE+	~8.5	~53.5
DAMO-YOLO	~9.5	~54.5
YOLOX	~10.5	~55.5
EfficientDet	~11.5	~56.5

[Documentation](#)

See below for quickstart installation and usage examples. For comprehensive guidance on training, validation, prediction, and deployment, refer to our full [Ultralytics Docs](#).

[README](#) [Code of conduct](#) [Apache-2.0 license](#) [Security](#)

## Latest updates -- SAM 2: Segment Anything in Images and Videos

Please check out our new release on [Segment Anything Model 2 \(SAM 2\)](#).

- SAM 2 code: <https://github.com/facebookresearch/segment-anything-2>
- SAM 2 demo: <https://sam2.metademolab.com/>
- SAM 2 paper: <https://arxiv.org/abs/2408.00714>

```

graph LR
    Image[Image] --> Encoder[Image encoder]
    Encoder --> Attention[Memory attention]
    Attention --> Decoder[Mask decoder]
    Decoder --> Mask[Mask]
    Decoder --> Encoder2[Memory encoder]
    Encoder2 --> Bank[Memory bank]
    Points[points] --> Decoder
    Box[box] --> Decoder
    Prompt[Prompt encoder] --> Decoder
    Mask --> Bank
    Bank --> Decoder
    
```

**Segment Anything Model 2 (SAM 2)** is a foundation model towards solving promptable visual segmentation in images and videos. We extend SAM to video by considering images as a video with a single frame. The model design is a simple transformer architecture with streaming memory for real-time video processing. We build a model-in-the-loop data engine, which improves model and data via user interaction, to collect [our SA-V dataset](#), the largest video segmentation dataset to date. SAM 2 trained on our data provides strong performance across a wide range of tasks and visual domains.

## Segment Anything

[Meta AI Research, FAIR](#)

Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, [Tete Xiao](#), Spencer Whitehead, Alex Berg, Wan-Yen Lo, Piotr Dollar, Ross Girshick

[[Paper](#)] [[Project](#)] [[Demo](#)] [[Dataset](#)] [[Blog](#)] [[BibTeX](#)]

```

graph LR
    Image[Image] --> Encoder[Image encoder]
    Encoder --> Embedding[Image embedding]
    Embedding --> Decoder[Mask decoder]
    Decoder --> Scissors[Scissors]
    Decoder --> Text[Text]
    Decoder --> Conv[Conv]
    Decoder --> Prompt[Prompt encoder]
    Conv --> Mask[Mask]
    Points[points] --> Mask
    Box[box] --> Mask
    Text --> Mask
    Mask --> Scissors
    
```

The **Segment Anything Model (SAM)** produces high quality object masks from input prompts such as points or boxes, and it can be used to generate masks for all objects in an image. It has been trained on a [dataset](#) of 11 million images and 1.1 billion masks, and has strong zero-shot performance on a variety of segmentation tasks.

Take home messages of the session:

- 1) CNNs are superior for cell detection and segmentation
- 2) Grad-CAM facilitates interpreting Deep Learning models
- 3) Lots of image analysis can be easily done via OpenCV
- 4) Mask-RCNN recommended for cell instance segmentation
- 5) YOLO, SAM and CellPose algorithms are becoming standard



# National Bioinformatics Infrastructure Sweden (NBIS)



*Knut och Alice  
Wallenbergs  
Stiftelse*



**LUNDS  
UNIVERSITET**